

Degree-2 FE (and More) from Bilinear Maps

1 Bilinear Maps

We introduce the notion of groups with bilinear maps. Let G_1 , G_2 and G_T be cyclic groups of order p , a prime, endowed with a map $e : G_1 \times G_2 \rightarrow G_T$ which obeys

- Non-degeneracy: If $g_1 \in G_1$ and $g_2 \in G_2$ are non-identity elements, $e(g_1, g_2)$ is a non-identity element as well.
- Efficient computability: There is a polynomial-time algorithm to compute e .
- Bilinearity: $e(g_1^{x_1}, g_2^{x_2}) = e(g_1, g_2)^{x_1 x_2}$. This lets you compute degree-2 polynomials “in the exponent”.

We can assume that the decisional Diffie-Hellman problem is hard in either G_1 , G_2 or G_T . That is,

$$(g, g^x, h, h^x) \approx_c (g, g^x, h, h^y)$$

for uniformly random $x, y \leftarrow \mathbb{Z}_p$ and a uniformly random elements $g, h \leftarrow G_1$ (or G_2 or G_T , respectively).

One has to be careful. If $g \leftarrow G_1$ and $h \leftarrow G_2$, this assumption is false. Can you see how?

Notation. We will let $[a]_1$ for $a \in \mathbb{Z}_p$ denote g_1^a for some fixed $g_1 \in G_1$. Similarly, $[a]_2$ for g_2^a with $g_2 \in G_2$, and $[a]_T$ for g_T^a for some $g_T \in G_T$.

2 Degree-1 FE: A Warmup

Consider a functional encryption for linear functions

$$F_{\alpha_1, \dots, \alpha_\ell}(x_1, \dots, x_\ell) = \sum_{i=1}^{\ell} \alpha_i x_i$$

We use an Abelian group G of order p , a prime.

- msk: $s_1, \dots, s_\ell \leftarrow \mathbb{Z}_p$ and mpk: $g^{s_1}, \dots, g^{s_\ell}$.
- Enc(mpk, (x_1, \dots, x_ℓ)): pick a random $r \leftarrow \mathbb{Z}_p$ and output

$$\text{ct} := (g^r, g^{r s_1 + x_1}, \dots, g^{r s_\ell + x_\ell})$$

(Note that one can compute $g^{r s_i + x_i}$ using the public key g^{s_i} alone; no need to know s_i .)

- Keygen(msk, $\alpha_1, \dots, \alpha_\ell$): outputs

$$\text{sk}_{\vec{\alpha}} = \sum_{i=1}^{\ell} \alpha_i s_i$$

- Dec: Compute

$$\prod_{i=1}^{\ell} g^{r s_i + x_i} / (g^r)^{\text{sk}_{\vec{\alpha}}} = g^{\sum_{i=1}^{\ell} \alpha_i x_i}$$

We will show security in just a moment, but a few things are worth noting.

1. We only get the output in the exponent. This is a shortcoming that will stay with us throughout this lecture (and indeed, something that we don't know how to overcome at all). Fortunately, at the end, in our application to XIO (cf. Neekon's lecture), we will compute functions whose outputs are binary, so given the output in the exponent, you can figure it out by brute force search.
2. There is no compression here. That is, by virtue of the function being linear, its description size (namely ℓ elements of \mathbb{Z}_p) is the same as the input length. Now, you could try using this to construct a degree-2 FE in the following manner. Write a quadratic function $Q(x_1, \dots, x_\ell) = \sum_{i,j} \alpha_{i,j} x_i x_j$ as a linear function

$$F_{\alpha}(\mathbf{x}) = \langle \alpha, \mathbf{x} \otimes \mathbf{x} \rangle$$

That is, a linear function of \mathbf{x} tensored with itself. However, now, the ciphertext size will be $O(\ell^2)$ making it non-compact. In our application to XIO via the JLS result, we will need the encryption to be compact, i.e. with size $O(\ell)$. To achieve this, we will need to turn to groups that admit bilinear maps.

2.1 Simulation Security

We will show how to simulate mpk, sk_f and ct_x given only f and $f(x)$. (The generalization to multiple functional secret keys follows analogously.) Roughly speaking, this shows that mpk, sk_f and ct_x together reveal no more than $f(x)$.

Let $f(x) = \sum_{i=1}^{\ell} \alpha_i x_i \bmod p$.

- Simulate the functional secret key sk_f for a linear function as a random group element s_f . Note that this is the correct distribution: the marginal distribution of sk_f is indeed uniformly random.
- Assume that $\alpha_1 \neq 0$. For fixed values of s_2, \dots, s_ℓ , note that

$$s_1 = (s_f - \sum_{i>1} \alpha_i s_i) / \alpha_1$$

- Simulate mpk as $(g^{s_1}, \dots, g^{s_\ell})$ where s_1 is computed as above. Note that this can be done using s_f and $g^{s_2}, \dots, g^{s_\ell}$.
- To simulate the ciphertext, note that

$$\begin{aligned} g^{rs_1+x_1} &= g^{rs_f \alpha_1^{-1}} \cdot \prod_{i>1} g^{-rs_i \alpha_i \alpha_1^{-1}} \cdot g^{x_1} \\ &= g^{rs_f \alpha_1^{-1}} \cdot \prod_{i>1} g^{-(rs_i+x_i) \alpha_i \alpha_1^{-1}} \cdot \prod_{i>1} g^{x_i \alpha_i \alpha_1^{-1}} \cdot g^{x_1} \\ &= g^{rs_f \alpha_1^{-1}} \cdot g^{(\sum \alpha_i x_i) \alpha_1^{-1}} \cdot \prod_{i>1} g^{-(rs_i+x_i) \alpha_i \alpha_1^{-1}} \end{aligned}$$

This can be generated using g^r , $g^{rs_1+x_1}$ and $\sum \alpha_i x_i$ alone (together with other public information the simulator knows, such as the α_i and s_f). In fact, $\sum \alpha_i x_i$ is unnecessary; it suffices to have $g^{\sum \alpha_i x_i}$, an observation that will come in handy later.

However, $g^r, g^{rs_2+x_2}, \dots, g^{rs_\ell+x_\ell}$ are pseudorandom by the Diffie-Hellman assumption. Thus, the simulator can do its job and generate a computationally indistinguishable simulation given only α_i and $\sum \alpha_i x_i$ (and in particular, no other information about the x_i).

2.2 An Extension

Before proceeding further, let us mention an extension of this scheme where the input \mathbf{x} and the function α are given in the exponent. That is, **Keygen** is given $[\alpha]_2$ and **Enc** is given $[\mathbf{x}]_1$.

- **msk**: $\mathbf{s} \leftarrow \mathbb{Z}_p^\ell$ and **mpk**: $[\mathbf{s}]_1$.
- **Enc**(**mpk**, \mathbf{x}): pick a random $r \leftarrow \mathbb{Z}_p$ and output

$$\text{ct} := [r]_1, [r\mathbf{s} + \mathbf{x}]_1$$

Note that encryption requires only $[\mathbf{x}]_1$, not \mathbf{x} “in the clear”.

- **Keygen**(**msk**, α): outputs

$$\text{sk}_\alpha = [\alpha^T \mathbf{s}]_2$$

Note that key generation requires only $[\alpha]_2$, not α “in the clear”.

- **Dec**: Compute, with the aid of the bilinear map,

$$[\alpha^T(r\mathbf{s} + \mathbf{x}) - r\alpha^T \mathbf{s}]_T = [\alpha^T \mathbf{x}]_T$$

Note that decryption also requires only $[\alpha]_2$, not α “in the clear”, and, as before, the result is obtained in the exponent of G_T .

We will leave the proof of security as an exercise but note that as above, simulation should only require $[\alpha^T \mathbf{x}]_1$ and not $\alpha^T \mathbf{x}$ in the clear.

2.3 FE Through The Lens of PSM Protocols

Modern FE constructions are best viewed from the lens of a simple, information-theoretic, game called a private simultaneous messages (or PSM) game. In such a game, Alice has a function f , Bob has an input x , and they share a private random string that we will suggestively call **msk**. The goal is for both of them to send a *single message* to Charlie at the end of which we want:

- **Correctness**: Charlie should be able to compute $f(x)$ given Alice and Bob’s messages.
- **Security**: Alice and Bob’s messages should reveal nothing other than $f(x)$. This is formalized via a simulator in the usual cryptographic way.

PSM for Linear Functions. Imagine a simple setting where Alice holds a vector α that defines a linear function, Bob holds an input \mathbf{x} and the function

$$F(\alpha, \mathbf{x}) = (\alpha, \langle \alpha, \mathbf{x} \rangle \bmod q)$$

A protocol to do this would have the shared random string be $\mathbf{s} \in \mathbb{Z}_p^n$, Bob sends $\mathbf{s} + \mathbf{x} \bmod p$ and Alice sends α and $\langle \alpha, \mathbf{s} \rangle \bmod p$. Charlie computes

$$\langle \alpha, \mathbf{s} + \mathbf{x} \rangle - \langle \alpha, \mathbf{s} \rangle \bmod p = \langle \alpha, \mathbf{x} \rangle \bmod p$$

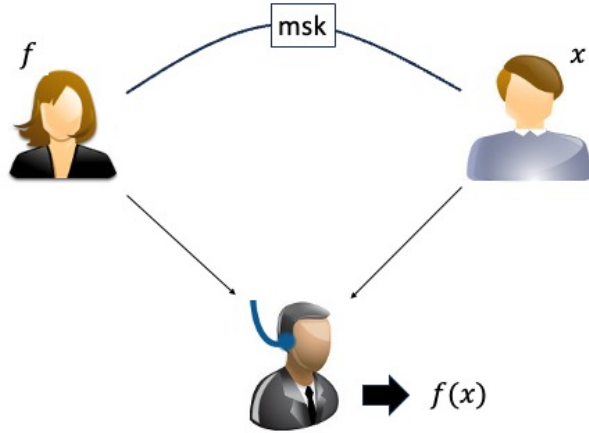


Figure 1: A Private Simultaneous Messages (PSM) Protocol.

as desired. In the language of FE, this gives us a one-ciphertext, one-key, secret-key FE scheme for linear functions.

It is clear that security is compromised if Bob uses the same msk to “encrypt” two different x . This is where groups and the Diffie-Hellman assumption come in and give us the scheme that we just presented. On the other hand, one can show that security is *not* compromised if Alice uses the same msk for many different linear functions. However, that is a fortuitous coincidence.

In general, the process of designing an FE scheme can be thought of as first designing a PSM protocol for the corresponding functionality, and then “lifting it” using cryptographic assumptions to many-key, many-ciphertext security, and possibly also to the public-key setting.

PSM for Inner Products. Imagine now that you want to do the same thing, but Alice wants to keep α private as well. Here is how one does this.

- msk is r, s .
- Alice sends $\alpha + r$ and Bob sends $x + s$.
- Charlie can now compute

$$\langle \alpha + r, x + s \rangle = \langle \alpha, x \rangle + \text{cross-terms}$$

where either Bob or Alice can compute each element of the cross-terms.

- In addition to the above, Alice sends Charlie $\langle \alpha, s \rangle + r'$ where msk contains a scalar r' in addition to the above.
- In addition to the above, Bob sends Charlie $\langle r, x + s \rangle - r'$.

I will leave checking correctness and proving security of this protocol as an exercise.

3 Degree-2 FE from Bilinear Maps

Let's start with a PSM protocol for the function

$$F(\alpha, (x, y)) = (\alpha, (x \otimes y) \cdot \alpha^T \bmod p) \quad (1)$$

where $x, y \in \mathbb{Z}_p^\ell$ and $\alpha \in \mathbb{Z}_p^{\ell^2}$. (Note that we treat vectors as row vectors so $x \otimes y$ is a row vector, and the column vector α^T denotes the transpose of α .)

We'd like the communication from Bob to be of size $O(\ell)$. As noted before, achieving communication $O(\ell^2)$ is trivial by appealing to the protocol for linear functions, but we want to (and need to) do better.

So, let's build up a PSM protocol for F from first principles.

- The msk is a pair of strings $r, s \leftarrow \mathbb{Z}_p^\ell$.
- Bob needs to hide x and y so he “one-time pads” them and sends $x + r \bmod p$ and $y + s \bmod p$ to Charlie.
- Alice sends α to Charlie.
- What can Charlie do with this information? He can compute

$$((x + r) \otimes (y + s)) \cdot \alpha^T = (x \otimes y) \cdot \alpha^T + \text{cross-terms} \quad (2)$$

so if he can remove the cross-terms somehow, he is all set.

Let's write down

$$\underbrace{((x + r) \otimes (y + s))}_{:=z} \cdot \alpha^T = (x \otimes y) \cdot \alpha^T + \underbrace{(x \otimes s) \cdot \alpha^T + (r \otimes z) \cdot \alpha^T}_{\text{cross-terms}} \quad (3)$$

There are two ways to handle the cross-terms. The first way is to observe that the cross terms are a linear function of x and z . Bob knows x and z , and Alice knows the description of the *secret* linear function (secret because it depends on r, s in addition to α .) Now, invoke the PSM protocol for computing inner products from our previous discussion. This will give us a PSM protocol for degree-2 functions with $O(\ell)$ communication.

The second way, which will lead us to our FE scheme, is to write down the cross-terms differently, as a *public* linear function of x, z, s, r . Now, the obvious way of doing this will result in a linear function in ℓ^2 dimensions and therefore a protocol with communication $O(\ell^2)$. But we will do better using cryptography.

Imagine for a moment that $r = ra_1$ and $s = sa_2$ are scalar multiples of *public* vectors a_1, a_2 . Then,

$$\text{cross-terms} = (x \otimes s) \cdot \alpha^T + (r \otimes z) \cdot \alpha^T = (x \otimes sa_2) \cdot \alpha^T + (ra_1 \otimes z) \cdot \alpha^T = [x \otimes s \parallel r \otimes z] \cdot \underbrace{\begin{bmatrix} I \otimes a_2 \\ a_1 \otimes I \end{bmatrix}}_{\text{linear fn}} \cdot \alpha^T \quad (4)$$

which is a public linear function in (sx, rz) . This can be computed using our first PSM protocol, the one for linear functions.

The problem, however, is that s and r are now not random and therefore leak information. We will solve this by “putting them in the exponent” and invoking an appropriate version of the Diffie-Hellman assumption.

3.1 The FE Scheme

A First Try. Our first attempt will have essentially all the ideas except it will be insecure because of the failure of the decisional Diffie-Hellman assumption. We will nevertheless be able to fix it under the k -LIN assumption for $k > 1$. (Recall that 1-LIN is DDH.)

- mpk is $[a_1]_1, [a_1]_2, [a_2]_2, [w]_1$, and $\text{msk} = (a_1, a_2, w)$.
- $\text{Enc}(\text{mpk}, (x, y))$ picks random $r, s \leftarrow \mathbb{Z}_p$ and computes $[x + ra_1]_1$ and $[y + sa_2]_2$. It also picks $s' \leftarrow \mathbb{Z}_p$ and computes $[s']_1$ and $[s'w + (sx\|rz)]_1$ and outputs all of them as the ciphertext. That is,

$$\text{ct} = ([x + ra_1]_1, [y + sa_2]_2, [s']_1, [s'w + (sx\|rz)]_1)$$

- $\text{Keygen}(\text{msk}, \alpha)$ outputs

$$\text{sk}_\alpha := [w\beta^T]_2$$

where

$$\beta^T = \begin{bmatrix} I \otimes a_2 \\ a_1 \otimes I \end{bmatrix} \cdot \alpha^T$$

Note that this can be computed given $[a_1]_2, [a_2]_2, w$ and α .

- Decryption uses the last two components of the ciphertext to compute

$$[(s'w + (sx\|rz)) \cdot \beta^T - s'(w\beta^T)]_T = [(sx\|rz) \cdot \beta^T]_T$$

using the fact that mpk allows Dec to compute $[\beta^T]_2$. One can now recover the output by computing

$$[(x + ra_1) \otimes (y + sa_2)] \cdot \alpha^T]_T = [(x \otimes y) \cdot \alpha^T + (sx\|rz) \cdot \beta^T]_T$$

and dividing by $[(sx\|rz) \cdot \beta^T]_T$.

All seems well except that given $[a_1]_2$, DDH does not hold in \mathbb{G}_1 w.r.t. a_1 , i.e. $[ra_1]_1$ does not look random any more. Indeed, letting the first coordinates of a_1 be (a_{11}, a_{12}) , we can check the equation

$$e([ra_{11}]_1, [a_{12}]_2) \stackrel{?}{=} e([a_{11}]_1, [ra_{12}]_2)$$

which is a distinguisher. Since $[ra_1]_1$ is not pseudorandom any more, we cannot use it as a one-time pad to hide x .

It seems necessary to publish $[a_1]_2$ in the mpk for the decryption algorithm to do its job. So, we seem to be stuck.

The Actual Scheme. Since DDH failed, let's lean in on k -LIN for $k > 1$. Recall that the DDH assumption says $[a]_1, [sa_1]_1$ is pseudorandom where $a \in \mathbb{Z}_p^{1 \times m}$ for a polynomially large m . We just saw that this is not true once one gets their hands on $[a]_2$ as well.

The (bilateral) k -LIN assumption says that

$$([A]_i, [sA]_i)_{i \in \{1,2,T\}} \approx_c ([A]_i, [u]_i)_{i \in \{1,2,T\}}$$

where $A \leftarrow \mathbb{Z}_p^{k \times m}$. Now, one can check that the attack we had before does not go through any more. This assumption is believed to hold and is a special case of an “uber-assumption” on bilinear maps [BBG05, Appendix A].

With this, one can write down the final scheme.

- mpk is $[A_1]_1, [A_1]_2, [a_2]_2, [w]_1$, and $\text{msk} = (A_1, a_2, w)$, where $A_1 \leftarrow \mathbb{Z}_p^{2 \times \ell}$.
- $\text{Enc}(\text{mpk}, (x, y))$ picks random $r \leftarrow \mathbb{Z}_p^2$, $s \leftarrow \mathbb{Z}_p$ and computes $[x + rA_1]_1$ and $[y + sa_2]_2$. It also picks $s' \leftarrow \mathbb{Z}_p$ and computes $[s']_1$ and $[s'w + (sx \parallel r \otimes z)]_1$ and outputs all of them as the ciphertext. That is,

$$\text{ct} = ([x + rA_1]_1, [y + sa_2]_2, [s']_1, [s'w + (sx \parallel r \otimes z)]_1)$$

- $\text{Keygen}(\text{msk}, \alpha)$ outputs

$$\text{sk}_\alpha := [w\beta^T]_2$$

where

$$\beta^T = \begin{bmatrix} I \otimes a_2 \\ A_1 \otimes I \end{bmatrix} \cdot \alpha^T$$

Note that this can be computed given $[A_1]_2, [a_2]_2, w$ and α .

- Decryption is exactly as before. It uses the last two components of the ciphertext to compute

$$[(s'w + (sx \parallel r \otimes z)) \cdot \beta^T - s'(w\beta^T)]_T = [(sx \parallel r \otimes z) \cdot \beta^T]_T$$

using the fact that mpk allows Dec to compute $[\beta^T]_2$. One can now recover the output by computing

$$[(x + rA_1) \otimes (y + sa_2)] \cdot \alpha^T]_T = [(x \otimes y) \cdot \alpha^T + (sx \parallel r \otimes z) \cdot \beta^T]_T$$

and dividing by $[(sx \parallel r \otimes z) \cdot \beta^T]_T$.

3.2 From Degree-(1, $O(1)$) FE to Degree-(2, $O(1)$) FE

We now need to compute

$$F(p, (\alpha, x, y)) = (p, \alpha, (x \otimes y) \cdot p(\alpha)^T \bmod p) \quad (5)$$

where $p(\alpha)$ denotes a vector each of whose coefficients is a degree- $O(1)$ polynomial computed on α . What we've been doing so far is the case of p being the constant function which ignores α and outputs a constant coefficient vector. Note that neither p nor α needs to be hidden.

One could go through the calculations as above and figure out that the cross terms are

$$\text{cross-terms} = [x \otimes s \parallel r \otimes z] \cdot \underbrace{\begin{bmatrix} I \otimes a \\ a \otimes I \end{bmatrix}}_{\text{linear fn}} \cdot p(\alpha)^T \quad (6)$$

which, as a function of x, y and α , is degree-1 in x and z and degree- $O(1)$ in α . Thus, given a degree-(1, $O(1)$)-FE scheme, we can immediately build a degree-(2, $O(1)$)-FE scheme.

We will not be able to get to the construction of a degree-(1, $O(1)$)-FE scheme in the lecture but (a) we will present a degree-(0, $O(1)$)-FE scheme, aka an ABE scheme for constant-degree polynomials. In fact, we will do one better and tell you about an ABE scheme for NC^1 functions, generalizing constant-degree polynomials; and (b) we will refer you to Wee's paper [Wee20] for the actual construction of degree-(1, $O(1)$)-FE.

4 ABE for NC¹ Functions

4.1 The GPSW ABE scheme

- $\text{msk} = s$ and $\text{mpk} = [t_1]_1, \dots, [t_\ell]_1, [s]_T$.
- $\text{Enc}(\text{mpk}, x_1, \dots, x_\ell, m)$ where $x_i \in \{0, 1\}$ computes

$$[rt_1 \cdot x_1]_1, \dots, [rt_\ell \cdot x_\ell]_1, [rs + m]_T$$

- $\text{Keygen}(\text{msk}, f)$ where f is an NC¹ circuit does the following: using the LSS scheme for f , compute a set of numbers $(\alpha_1, \dots, \alpha_m)$ together with a mapping $\rho : [m] \rightarrow [\ell]$ denoting which index / party each share belongs to. The secret key for f consists of

$$([\alpha_i \cdot t_{\rho(i)}^{-1}]_2)_{i \in [m]}$$

- Dec first computes

$$[r\alpha_i]_T$$

for each $i \in [m]$ for which $x_{\rho(i)} = 1$. From this, using LSS reconstruction, compute $[rs]_T$ and therefore $[m]_T$.

References

- [BBG05] Dan Boneh, Xavier Boyen, and Eu-Jin Goh. Hierarchical identity based encryption with constant size ciphertext. Cryptology ePrint Archive, Paper 2005/015, 2005.
- [Wee20] Hoeteck Wee. Functional encryption for quadratic functions from k-lin, revisited. In *Theory of Cryptography: 18th International Conference, TCC 2020, Durham, NC, USA, November 16–19, 2020, Proceedings, Part I* 18, pages 210–228. Springer, 2020.